

```
In [24]: 1 import numpy as np
          2 import pandas as pd
          3 import seaborn as sns
          4 import matplotlib.pyplot as plt
          5 from sklearn.linear_model import LinearRegression
          6 from sklearn.linear_model import Ridge,RidgeCV,Lasso
          7 from sklearn.preprocessing import StandardScaler
          8 from sklearn.model_selection import train_test_split
```

```
In [25]: 1 df=pd.read_csv(r"C:\Users\yoshitha lakshmi\OneDrive\Desktop\python\Advertising.csv")
          2 df
```

```
Out[25]:
```

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9
...	...	...	...	...
195	38.2	3.7	13.8	7.6
196	94.2	4.9	8.1	14.0
197	177.0	9.3	6.4	14.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	18.4

200 rows × 4 columns

In [26]:

```
1 df.head()
```

Out[26]:

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9

In [27]:

```
1 df.tail()
```

Out[27]:

	TV	Radio	Newspaper	Sales
195	38.2	3.7	13.8	7.6
196	94.2	4.9	8.1	14.0
197	177.0	9.3	6.4	14.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	18.4

In [28]:

```
1 df.shape
```

Out[28]: (200, 4)

In [29]:

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0    TV          200 non-null    float64
1    Radio        200 non-null    float64
2    Newspaper    200 non-null    float64
3    Sales        200 non-null    float64
dtypes: float64(4)
memory usage: 6.4 KB
```

In [30]:

```
1 df.describe()
```

Out[30]:

	TV	Radio	Newspaper	Sales
<b>count</b>	200.000000	200.000000	200.000000	200.000000
<b>mean</b>	147.042500	23.264000	30.554000	15.130500
<b>std</b>	85.854236	14.846809	21.778621	5.283892
<b>min</b>	0.700000	0.000000	0.300000	1.600000
<b>25%</b>	74.375000	9.975000	12.750000	11.000000
<b>50%</b>	149.750000	22.900000	25.750000	16.000000
<b>75%</b>	218.825000	36.525000	45.100000	19.050000
<b>max</b>	296.400000	49.600000	114.000000	27.000000

In [31]:


```
1 features = ['TV', 'Radio', 'Newspaper']
```

In [32]:

```
1 target = ['Sales']
```

```
In [33]: 1 sns.pairplot(df,x_vars=['TV','Radio','Newspaper'],y_vars=['Sales'],height=7,aspect=0.7,kind='reg')
```

```
Out[33]: <seaborn.axisgrid.PairGrid at 0x1acdfea2080>
```



```
In [34]: 1 from sklearn.model_selection import train_test_split  
2 from sklearn.linear_model import LinearRegression
```


```
In [55]: 1 x=np.array(df[features])  
2 y=np.array(df[target])
```

```
In [66]: 1 X_train,X_test,y_train,y_test = train_test_split(x,y,test_size=0.25)
          2 regr = LinearRegression()
          3 regr.fit(X_train,y_train)
          4 print(regr.score(x_test, y_test))
```

-0.5824704202324338

```
In [67]: 1 sns.heatmap(df.corr())
```

Out[67]: <Axes: >



```
In [68]: 1 from sklearn.linear_model import LinearRegression
          2 lm = LinearRegression()
          3 lm.fit(X_train,y_train)
```

Out[68]: LinearRegression()

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [69]: 1 print(lm.intercept_)
```

4.660916410008371

```
In [70]: 1 df.columns
```

Out[70]: Index(['TV', 'Radio', 'Newspaper', 'Sales'], dtype='object')

```
In [71]: 1 Addf=df[['TV', 'Radio', 'Newspaper', 'Sales']]
```

```
In [72]: 1 X = Addf[['TV', 'Radio', 'Newspaper']]
          2 y = df['Sales']
```

```
In [73]: 1 from sklearn.linear_model import LinearRegression
          2 lm = LinearRegression()
          3 lm.fit(X_train,y_train)
```

Out[73]: LinearRegression()

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [74]: 1 coeff_df=pd.DataFrame(lm.coef_,X.columns,columns=['coefficient'])  
        2 coeff_df
```

```
Out[74]:
```

	coefficient
TV	0.054360
Radio	0.108862
Newspaper	-0.002042

```
In [75]: 1 predictions=lm.predict(x_test)
```

```
In [76]: 1 from sklearn import metrics  
        2 print('MAE:',metrics.mean_absolute_error(y_test,predictions))  
        3 print('MSE:',metrics.mean_squared_error(y_test,predictions))  
        4 print('RMSE:',np.sqrt(metrics.mean_squared_error(y_test,predictions)))
```

MAE: 5.609455830624045

MSE: 47.55186887354155

RMSE: 6.8957863129262895

## Ridge Regression Model

```
In [77]: 1 ridgeReg = Ridge(alpha=10)
          2 ridgeReg.fit(X_train,y_train)
          3 train_score_ridge = ridgeReg.score(X_train,y_train)
          4 test_score_ridge = ridgeReg.score(x_test,y_test)
          5
          6 print('\nRidge Model\n')
          7 print('Train score for ridge model is {}'.format(train_score_ridge))
          8 print('Test score for ridge model is{}'.format(test_score_ridge))
```

Ridge Model

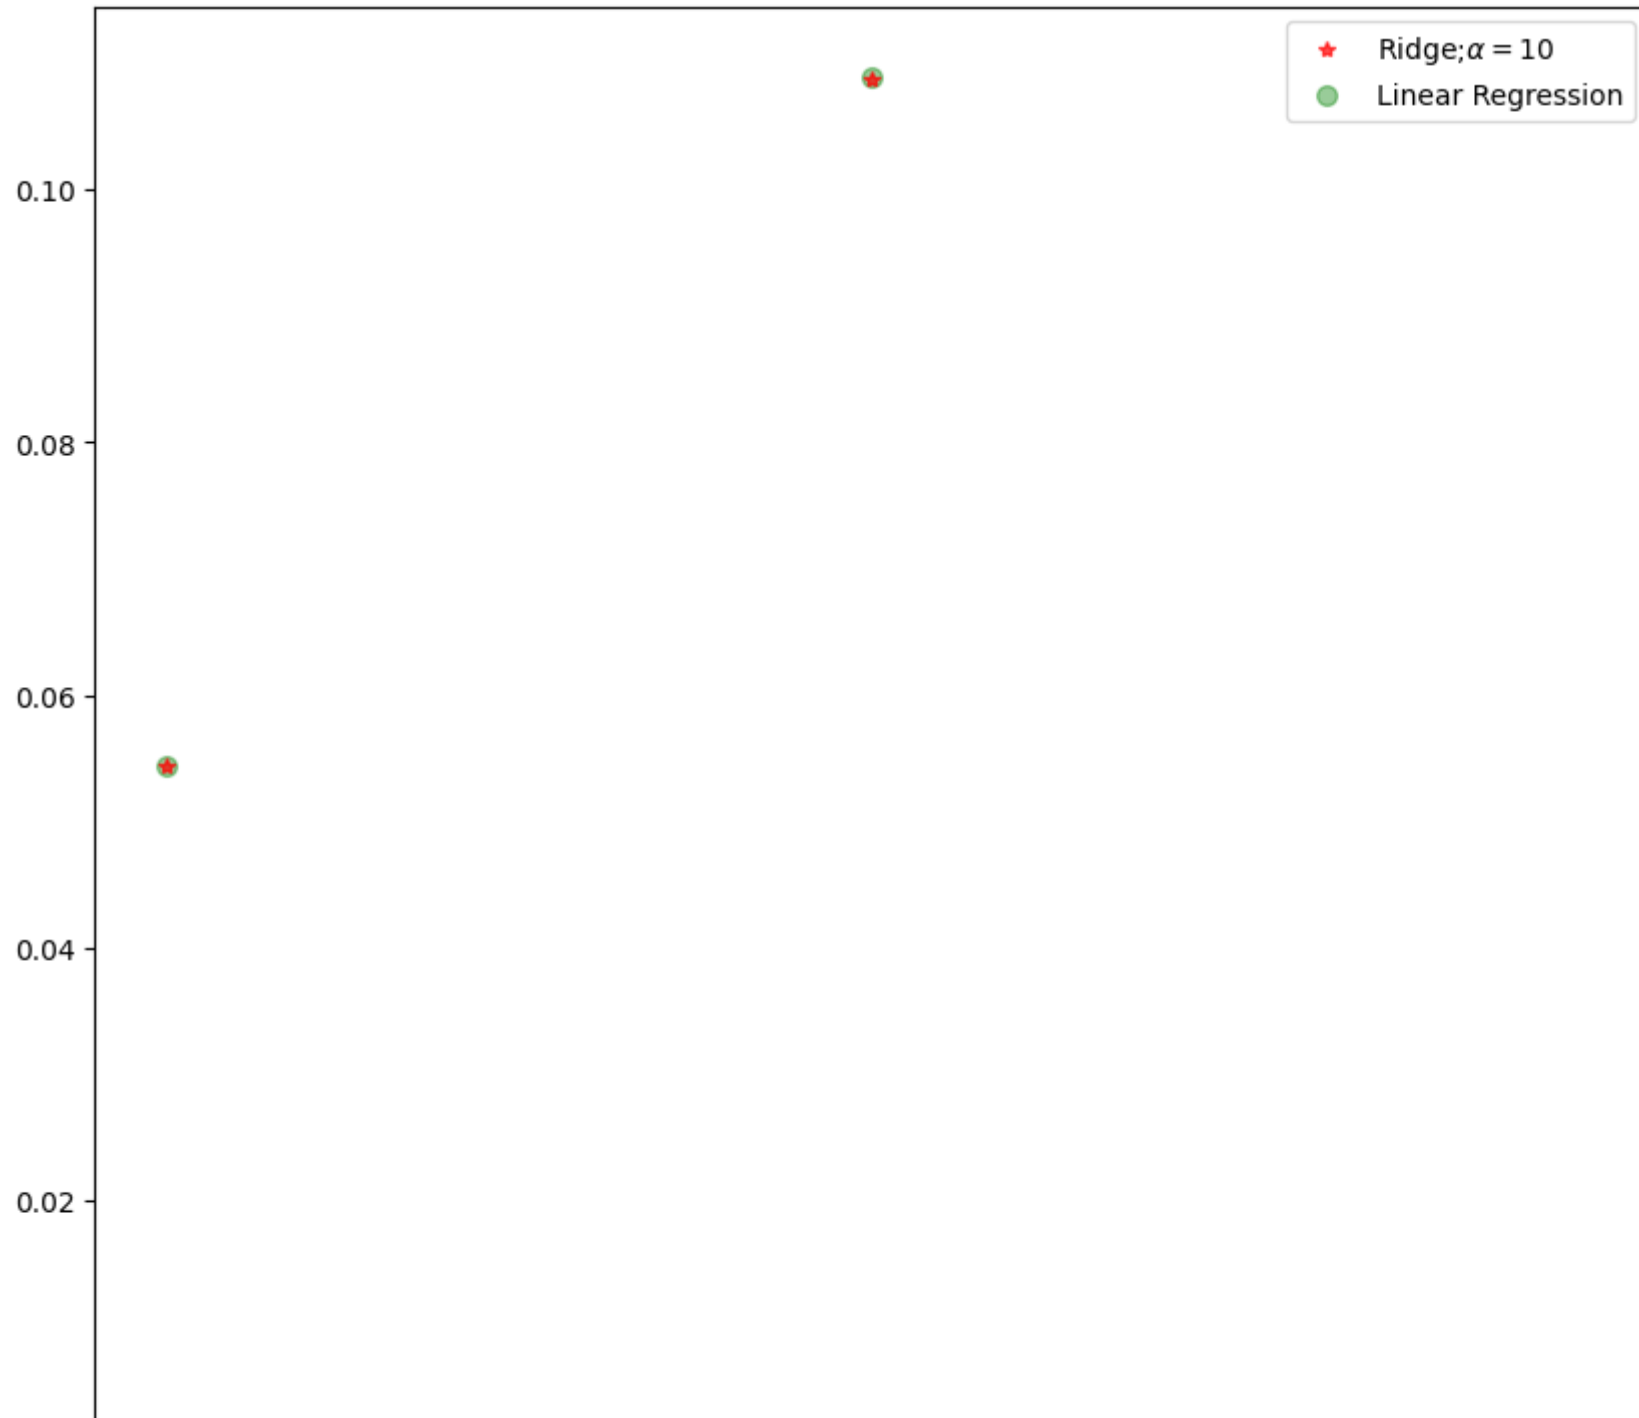
Train score for ridge model is 0.8954000079149985


Test score for ridge model is-0.5824696427173317



```
In [78]: 1 plt.figure(figsize = (10,10))
2 plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=6,color='red',label=r'Ridge;$\alpha=0.7$')
3 plt.plot(features,lm.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='Linear Regres')
4 plt.xticks(rotation=90)
5 plt.legend()
6 plt.show()
```







```
In [79]: 1 lassoReg = Lasso(alpha=10)
2 lassoReg.fit(X_train,y_train)
3 train_score_lasso = lassoReg.score(X_train,y_train)
4 test_score_lasso = lassoReg.score(x_test,y_test)
5
6 print('\nRidge Model\n')
7 print('Train score for lasso model is {}'.format(train_score_lasso))
8 print('Test score for lasso model is{}'.format(test_score_lasso))
```


Ridge Model

Train score for lasso model is 0.8789640031484611  
Test score for lasso model is-0.5733638138042045

```
In [80]: 1 plt.figure(figsize = (10,10))
2 plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=6,color='red',label=r'Lasso;$\lambda$')
3 plt.plot(features,lm.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='Linear Regres')
4 plt.xticks(rotation=90)
5 plt.legend()
6 plt.show()
```











In [81]:

```
1 # Comparison between Ridge,Lasso and RidgeCV
2 plt.figure(figsize=(10,10))
3
4 plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=6,color='red',label=r'Ridge;$\
5 plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='d',markersize=6,color='red',label=r'Lasso;$\
6 plt.plot(features,lm.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='Linear Regres
7 plt.xticks(rotation=90)
8 plt.title('comparison between ridge,lasso and linear')
9 plt.legend()
10 plt.show()
```







```
In [82]: 1 # Linear CV model using Ridge
2 from sklearn.linear_model import RidgeCV
3 ridge_CV=RidgeCV(alphas=[0.0001,0.01,0.001,0.1]).fit(X_train,y_train)
4 print("The train score for ridge model is {}".format(ridge_CV.score(X_train,y_train)))
5 print("The test score for ridge model is {}".format(ridge_CV.score(x_test,y_test)))
```

The train score for ridge model is 0.8954000179269546

The test score for ridge model is -0.5824704108926089

```
In [83]: 1 # Linear CV model using Lasso
2 from sklearn.linear_model import LassoCV
3 lasso_CV=LassoCV(alphas=[0.0001,0.01,0.001,0.1]).fit(X_train,y_train)
4 print("The train score for lasso model is {}".format(lasso_CV.score(X_train,y_train)))
5 print("The test score for lasso model is {}".format(lasso_CV.score(x_test,y_test)))
```

The train score for lasso model is 0.8954000179239859

The test score for lasso model is -0.5824703283838824

## Elastic Net

```
In [87]: 1 from sklearn.linear_model import ElasticNet
2 regr = ElasticNet()
3 regr.fit(x,y)
4 print(regr.coef_)
5 print(regr.intercept_)
6 y_pred_elastic = regr.predict(X_train)
7 mean_squared_error = np.mean((y_pred_elastic-y_train)**2)
8 print('Mean squared error on test set',mean_squared_error)
```

```
[0.05440081 0.1046715 0.          ]
4.696191158087226
Mean squared error on test set 2.8150966541195057
```

```
In [ ]:
```

```
1
```